

```
In [1]: import os
import pandas as pd
```

```
In [2]: current_directory = os.getcwd()
print("Current Directory:", current_directory)
```

Current Directory: /Users/amarachiordor/Documents/Portfolio projects

```
In [3]: files = os.listdir()
print("Files in Directory:", files)
```

Files in Directory: ['Credic card fraud detection main.ipynb', 'creditcard.csv', 'Creditcard fraud detection.ipynb', '.ipynb_checkpoints']

```
In [4]: #Load CSV File
df = pd.read_csv(r"/Users/amarachiordor/Documents/Portfolio projects/cred
```

```
In [5]: df.head()
```

```
Out[5]:
```

	Time	V1	V2	V3	V4	V5	V6	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.2395
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.0788
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.7914
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.2376
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.5929

5 rows × 31 columns

```
In [6]: pd.options.display.max_columns = None
```

```
In [7]: df.shape
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Time        284807 non-null float64
 1   V1           284807 non-null float64
 2   V2           284807 non-null float64
 3   V3           284807 non-null float64
 4   V4           284807 non-null float64
 5   V5           284807 non-null float64
 6   V6           284807 non-null float64
 7   V7           284807 non-null float64
 8   V8           284807 non-null float64
 9   V9           284807 non-null float64
10  V10          284807 non-null float64
11  V11          284807 non-null float64
12  V12          284807 non-null float64
13  V13          284807 non-null float64
14  V14          284807 non-null float64
15  V15          284807 non-null float64
16  V16          284807 non-null float64
17  V17          284807 non-null float64
18  V18          284807 non-null float64
19  V19          284807 non-null float64
20  V20          284807 non-null float64
21  V21          284807 non-null float64
22  V22          284807 non-null float64
23  V23          284807 non-null float64
24  V24          284807 non-null float64
25  V25          284807 non-null float64
26  V26          284807 non-null float64
27  V27          284807 non-null float64
28  V28          284807 non-null float64
29  Amount       284807 non-null float64
30  Class        284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

```
In [8]: df.shape
```

```
Out[8]: (284807, 31)
```

```
In [9]: #Checking for missing values
df.isnull().sum()
```

```
Out [9]: Time      0
         V1        0
         V2        0
         V3        0
         V4        0
         V5        0
         V6        0
         V7        0
         V8        0
         V9        0
         V10       0
         V11       0
         V12       0
         V13       0
         V14       0
         V15       0
         V16       0
         V17       0
         V18       0
         V19       0
         V20       0
         V21       0
         V22       0
         V23       0
         V24       0
         V25       0
         V26       0
         V27       0
         V28       0
         Amount    0
         Class     0
         dtype: int64
```

```
In [10]: from sklearn.preprocessing import StandardScaler
```

```
In [11]: sc = StandardScaler()
         df['Amount'] = sc.fit_transform(pd.DataFrame(df['Amount']))
```

```
In [12]: df = df.drop(['Time'],axis = 1)
```

```
In [13]: # Check for Duplicates
         df.duplicated().sum()
```

```
Out [13]: 9144
```

```
In [14]: df = df.drop_duplicates()
```

```
In [15]: df.shape
```

```
Out [15]: (275663, 30)
```

```
In [16]: df['Class'].value_counts()
```

```
Out [16]: Class
         0      275190
         1        473
         Name: count, dtype: int64
```

```

In [17]: X = df.drop ('Class',axis = 1)
         y = df['Class']

In [18]: from sklearn.model_selection import train_test_split

In [19]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2

In [20]: import numpy as np
         from sklearn.linear_model import LogisticRegression
         from sklearn.ensemble import RandomForestClassifier
         from sklearn.tree import DecisionTreeClassifier
         from sklearn.metrics import accuracy_score, f1_score, precision_score, re

In [21]: classifier = {
         "Logistic Regression": LogisticRegression(),
         "Decision Tree Classifier": DecisionTreeClassifier ()
         }

         for name, clf in classifier.items():
             print (f"\n===== {name} =====")
             clf.fit(X_train, y_train)
             y_pred = clf.predict(X_test)
             print(f"\n Accuracy: {accuracy_score(y_test, y_pred)} ")
             print(f"\n Precision: {precision_score(y_test, y_pred)} ")
             print(f"\n Recall: {recall_score(y_test, y_pred)} ")
             print(f"\n F1 score: {f1_score(y_test, y_pred)} ")

=====Logistic Regression=====

Accuracy: 0.9992563437505668

Precision: 0.890625

Recall: 0.6263736263736264

F1 score: 0.7354838709677419

=====Decision Tree Classifier=====

Accuracy: 0.998911722561805

Precision: 0.6631578947368421

Recall: 0.6923076923076923

F1 score: 0.6774193548387096

In [ ]: #Under sampling

In [45]: normal = df[df['Class']==0]
         fraud = df[df['Class']==1]

In [47]: normal.shape

Out[47]: (275190, 30)

In [49]: fraud.shape

```

Out[49]: (473, 30)

In [51]: `normal_sample = normal.sample(n = 473)`

In [53]: `normal_sample.shape`

Out[53]: (473, 30)

In [55]: `new_df = pd.concat([normal_sample, fraud], ignore_index = True)`

In [57]: `new_df.head()`

Out[57]:

	V1	V2	V3	V4	V5	V6	V7
0	-0.734589	0.746812	1.924888	-1.248523	0.453207	-0.143757	0.406034
1	0.652989	-1.138956	0.853234	0.052937	-0.473422	2.044374	-0.847534
2	-0.756901	1.523053	-2.686048	-0.738390	0.294788	-0.640744	0.139390
3	-0.804967	1.128930	-1.057995	-1.930522	1.830317	4.137497	-2.143935
4	1.493342	-0.980063	0.570087	-1.237290	-1.762276	-1.231630	-0.936145

In [59]: `new_df['Class'].value_counts()`

Out[59]: Class
 0 473
 1 473
 Name: count, dtype: int64

In [61]: `X = new_df.drop('Class', axis = 1)`
`y = new_df['Class']`

In [63]: `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)`

In [65]:

```

classifier = {
    "Logistic Regression": LogisticRegression(),
    "Decision Tree Classifier": DecisionTreeClassifier ()
}

for name, clf in classifier.items():
    print (f"\n===== {name} =====")
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    print(f"\n Accuracy: {accuracy_score(y_test, y_pred)} ")
    print(f"\n Precision: {precision_score(y_test, y_pred)} ")
    print(f"\n Recall: {recall_score(y_test, y_pred)} ")
    print(f"\n F1 score: {f1_score(y_test, y_pred)} ")

```

=====Logistic Regression=====

Accuracy: 0.9526315789473684

Precision: 0.979381443298969

Recall: 0.9313725490196079

F1 score: 0.9547738693467337

=====Decision Tree Classifier=====

Accuracy: 0.9210526315789473

Precision: 0.9393939393939394

Recall: 0.9117647058823529

F1 score: 0.9253731343283582

In [67]: `#Oversampling`

In [69]: `X = df.drop ('Class',axis = 1)
y = df['Class']`

In [71]: `X.shape`

Out[71]: (275663, 29)

In [73]: `y.shape`

Out[73]: (275663,)

In [75]: `from imblearn.over_sampling import SMOTE`

In [77]: `x_res, y_res = SMOTE().fit_resample(X,y)`

In [79]: `y_res.value_counts()`

Out[79]: Class
0 275190
1 275190
Name: count, dtype: int64

In [81]: `X_train, X_test, y_train, y_test = train_test_split(x_res, y_res, test_si`

In [83]: `classifier = {
 "Logistic Regression": LogisticRegression(),
 "Decision Tree Classifier": DecisionTreeClassifier ()
}

for name, clf in classifier.items():
 print (f"\n===== {name} =====")
 clf.fit(X_train, y_train)
 y_pred = clf.predict(X_test)
 print(f"\n Accuracy: {accuracy_score(y_test, y_pred)} ")
 print(f"\n Precision: {precision_score(y_test, y_pred)} ")`

```
print(f"\n Recall: {recall_score(y_test, y_pred)} ")
print(f"\n F1 score: {f1_score(y_test, y_pred)} ")
```

=====Logistic Regression=====

Accuracy: 0.9443657109633344

Precision: 0.9730561523721039

Recall: 0.9139683290002364

F1 score: 0.9425871411696323

=====Decision Tree Classifier=====

Accuracy: 0.9982284966750246

Precision: 0.997458611675864

Recall: 0.9990000545424795

F1 score: 0.9982287380439818

```
In [89]: dtc = DecisionTreeClassifier()
         dtc.fit(x_res, y_res)
```

```
Out[89]: ▼ DecisionTreeClassifier ⓘ ⓘ
         DecisionTreeClassifier()
```

```
In [90]: import joblib
```

```
In [95]: joblib.dump(dtc, "credit_card_model.pkl")
```

```
Out[95]: ['credit_card_model.pkl']
```

```
In [97]: model = joblib.load("credit_card_model.pkl")
```

```
In [101... pred = model.predict([[ 1.22965763450793,      0.141003507049326,      0
```

```
/opt/anaconda3/lib/python3.12/site-packages/sklearn/base.py:493: UserWarni
ng: X does not have valid feature names, but DecisionTreeClassifier was fi
tted with feature names
  warnings.warn(
```

```
In [103... pred[0]
```

```
Out[103... 0
```

```
In [105... if pred == 0:
            print("Normal Transaction")
        else:
            print("Fraud Transaction")
```

Normal Transaction

```
In [ ]:
```